

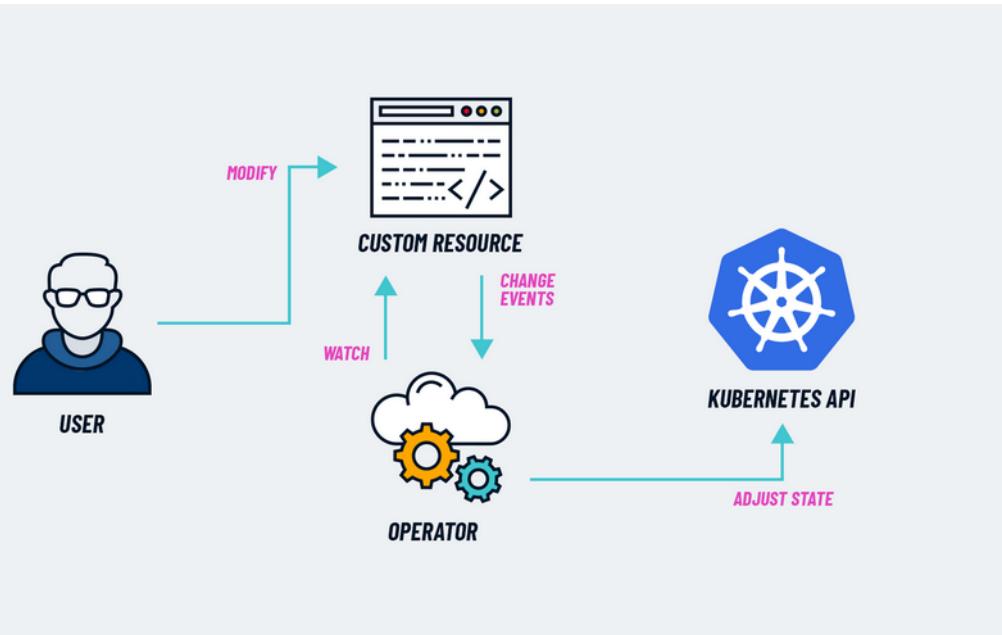
Kubernetes

mardi 7 février 2023 09:27

<https://kubernetes.io/fr/>

<https://fr.wikipedia.org/wiki/Kubernetes>

<https://wiki.archlinux.org/title/Kubernetes>

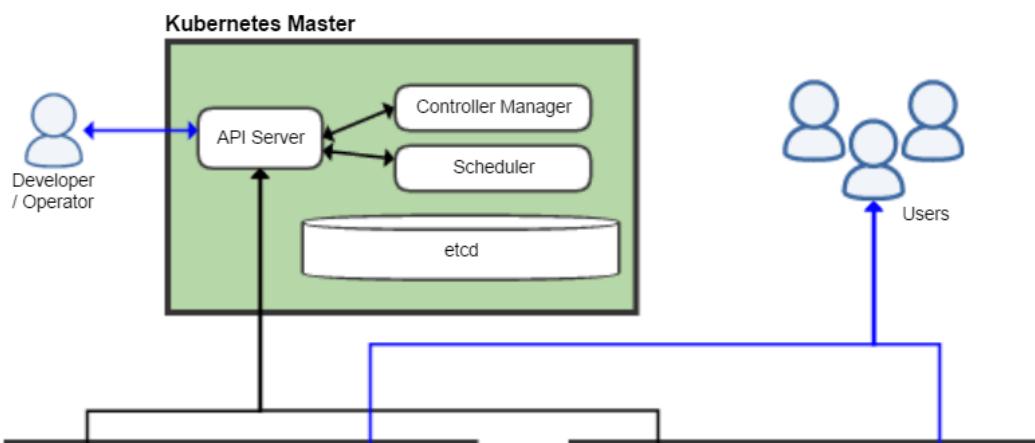


Kubernetes

[Article](#) [Discussion](#)

Kubernetes (communément appelé « K8s² ») est un système [open source](#) qui vise à fournir une « [plate-forme](#) permettant d'automatiser le déploiement, la montée en charge et la mise en œuvre de [conteneurs d'application](#) sur des [grappes de serveurs³](#) ». Il fonctionne avec toute une série de technologies de conteneurisation, et est souvent utilisé avec [Docker](#). Il a été conçu à l'origine par [Google](#), puis offert à la [Cloud Native Computing Foundation](#).

Dans le principe, Kubernetes reprend les mêmes bases qu'Ansible : déployer des applications et les gérer



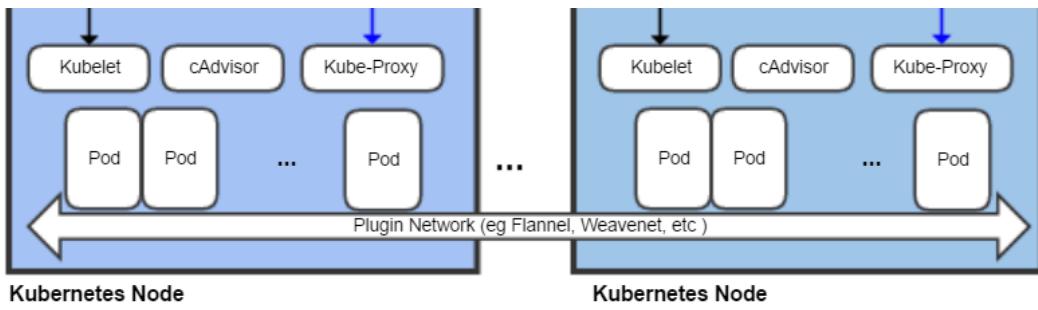
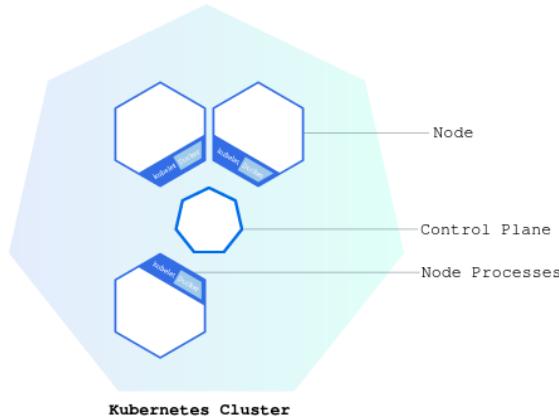


Schéma du Cluster



Un pod représente une ou plusieurs application (un pod pourrait contenir le front et un autre le back par exemple)

Pods [modifier | modifier le code]

L'unité de base de l'ordonnancement dans Kubernetes est appelée « *pod* ». C'est une vue abstraite de composants conteneurisés. Un *pod* consiste en un ou plusieurs conteneurs qui ont la garantie d'être co-localisés sur une machine hôte et peuvent en partager les ressources¹⁴. Chaque *pod* dans Kubernetes possède une [adresse IP](#) unique (à l'intérieur du cluster), qui permet aux applications d'utiliser les ports de la machine sans risque de conflit¹⁵. Un *pod* peut définir un volume, comme un répertoire sur un disque local ou sur le réseau, et l'exposer aux conteneurs de ce *pod*¹⁶. Les *pods* peuvent être gérés manuellement au travers de l'[API](#) de Kubernetes. Leur gestion peut également être déléguée à un contrôleur¹⁴.

Les *pods* sont rattachés au nœud qui les déploie jusqu'à leur expiration ou leur suppression. Si le nœud est défaillant, de nouveaux *pods* possédant les mêmes propriétés que les précédents seront déployés sur d'autres nœuds disponibles¹⁷.

Tutoriel :

<https://kubernetes.io/fr/docs/tutorials/>

Utiliser Minikube pour créer un cluster : <https://kubernetes.io/fr/docs/tutorials/kubernetes-basics/create-cluster/cluster-intro/>

Objectifs

- Découvrez ce qu'est un cluster Kubernetes.
- Apprenez ce qu'est Minikube.
- Démarrez un cluster Kubernetes à l'aide d'un terminal en ligne.

Un cluster Kubernetes est constitué de deux types de ressources:

- Le **maître** coordonne le cluster.
- Les **nœuds** sont les serveurs qui exécutent des applications.

Cluster up and running

We already installed minikube for you. Check that it is properly installed, by running the *minikube version* command:

```
minikube version ↵
```

OK, we can see that minikube is in place.

Start the cluster, by running the *minikube start* command:

```
minikube start ↵
```

Great! You now have a running Kubernetes cluster in your online terminal. Minikube started a virtual machine for you, and a Kubernetes cluster is now running in that VM.

```
$ minikube version
minikube version: v1.18.0
commit: ec61815d60f66a6e4f6353030a40b12362557caa-dirty
$ minikube start
* minikube v1.18.0 on Ubuntu 18.04 (amd64)
* Using the none driver based on existing profile

X The requested memory allocation of 2200MiB does not leave room for system overheat (total system memory: 2460MiB). You may face stability issues.
* Suggestion: Start minikube with less memory allocated: 'minikube start --memory=2200mb'

* Starting control plane node minikube in cluster minikube
* Running on localhost (CPUs=2, Memory=2460MB, Disk=194868MB) ...
* OS release is Ubuntu 18.04.5 LTS
* Preparing Kubernetes v1.20.2 on Docker 19.03.13 ...
  - kubelet.resolv-conf=/run/systemd/resolve/resolv.conf
  - Generating certificates and keys ...
  - Booting up control plane ...
  - Configuring RBAC rules ...
* Configuring local host environment ...
* Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v4
* Enabled addons: default-storageclass, storage-provisioner
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
$ []
```

Cluster version

To interact with Kubernetes during this bootcamp we'll use the command line interface, `kubectl`.

We'll explain `kubectl` in detail in the next modules, but for now, we're just going to look at some cluster information. To check if `kubectl` is installed you can run the *kubectl version* command:

```
kubectl version ↵
```

```
$ kubectl version
Client Version: version.Info{Major:"1", Minor:"20", GitVersion:"v1.20.4", GitCommit:"e87da0bd6e03ec3fea7933c4b5263d151aa07c", GitTreeState:"clean", BuildDate:"2021-02-18T16:12:00Z", GoVersion:"go1.15.8", Compiler:"gc", Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"20", GitVersion:"v1.20.2", GitCommit:"faecb196815e248d3ecfb03c680a4507229c2a56", GitTreeState:"clean", BuildDate:"2021-01-13T13:20:00Z", GoVersion:"go1.15.5", Compiler:"gc", Platform:"linux/amd64"}
$ []
```

Cluster details

Let's view the cluster details. We'll do that by running `kubectl cluster-info`:

```
kubectl cluster-info ↵
```

During this tutorial, we'll be focusing on the command line for deploying and exploring our application. To view the nodes in the cluster, run the `kubectl get nodes` command:

```
kubectl get nodes ↵
```

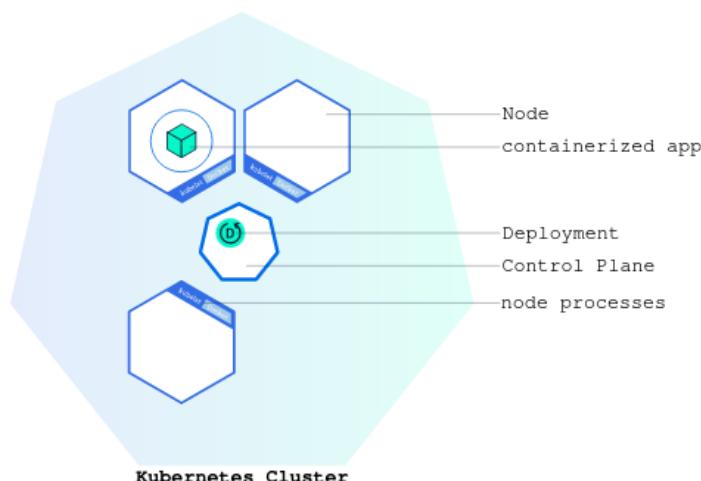
This command shows all nodes that can be used to host our applications. Now we have only one node, and we can see that its status is ready (it is ready to accept applications for deployment).

```
Kubernetes Bootcamp Terminal

$ minikube version
minikube version: v1.18.0
commit: ec61815d60f66a6e4f6353030a40b12362557caa-dirty
$ kubectl cluster-info
Kubernetes control plane is running at https://10.0.0.81:8443
KubeDNS is running at https://10.0.0.81:8443/api/v1/namespaces/kube-system/service/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
$ kubectl get nodes
NAME      STATUS    ROLES          AGE     VERSION
minikube  Ready     control-plane,master  3m7s   v1.20.2
$
```

Deploying your first app on Kubernetes



< | Module 2 - Deploy an app 58:32

Step 1 of 3 ▶

KUBECTL BASICS

Like minikube, kubectl comes installed in the online terminal. Type kubectl in the terminal to see its usage. The common format of a kubectl command is: kubectl action resource. This performs the specified action (like create, describe) on the specified resource (like node, container). You can use `--help` after the command to get additional info about possible parameters (`kubectl get nodes --help`).

Check that kubectl is configured to talk to your cluster, by running the `kubectl version` command:

```
kubectl version ↵
```

```
$ sleep 1; launch.sh
Starting Kubernetes. This is expected to take less than a minute.....
Kubernetes Started
$
$
$ kubectl version
Client Version: version.Info{Major:"1", Minor:"20", GitVersion:"v1.20.4", GitCommit:"e87da0bd6e03ec3fea7933c4b5263d151aaf07c", GitTreeState:"clean", BuildDate:"2021-02-18T16:12:00Z", GoVersion:"go1.15.8", Compiler:"gc", Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"20", GitVersion:"v1.20.2", GitCommit:"faecb196815e248d3ecfb03c680a4507229c2a56", GitTreeState:"clean", BuildDate:"2021-01-13T13:20:00Z", GoVersion:"go1.15.5", Compiler:"gc", Platform:"linux/amd64"}
$ ↵
```

OK, kubectl is installed and you can see both the client and the server versions.

To view the nodes in the cluster, run the `kubectl get nodes` command:

```
kubectl get nodes ↵
```

Here we see the available nodes (1 in our case). Kubernetes will choose where to deploy our application based on Node available resources.

```
$ kubectl get nodes
NAME      STATUS    ROLES          AGE      VERSION
minikube  Ready     control-plane,master  3m23s   v1.20.2
$ ↵
```

Deploy our app

Let's deploy our first app on Kubernetes with the `kubectl create deployment` command. We need to provide the deployment name and app image location (include the full repository url for

image location (include the full repository URL for images hosted outside Docker hub).

```
kubectl create deployment kubernetes-bootcamp --image=gcr.io/google-samples/kubernetes-bootcamp:v1 ✓
```

```
minikube ready control-plane,master   3m23s   v1.20.2
$ kubectl create deployment kubernetes-bootcamp --image=gcr.io/google-samples/kubernetes-bootcamp:v1
deployment.apps/kubernetes-bootcamp created
```

Great! You just deployed your first application by creating a deployment. This performed a few things for you:

- searched for a suitable node where an instance of the application could be run (we have only 1 available node)
- scheduled the application to run on that Node
- configured the cluster to reschedule the instance on a new Node when needed

To list your deployments use the

```
get deployments
```

```
kubectl get deployments ↵
```

```
$ kubectl get deployments
NAME                  READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp   1/1     1           1           2m19s
$ ↴
```

View our app

Pods that are running inside Kubernetes are running on a private, isolated network. By default they are visible from other pods and services within the same Kubernetes cluster, but not outside that network. When we use `kubectl`, we're interacting through an API endpoint to communicate with our application.

We will cover other options on how to expose your application outside the Kubernetes cluster in Module 4.

The `kubectl` command can create a proxy that will forward communications into the cluster-wide, private network. The proxy can be terminated by pressing control-C and won't show any output while it's running.

We will open a second terminal window to run the proxy.

```
echo -e "\n\n\n\033[92mStarting Proxy.\nAfter starting it will not output a\nresponse. Please click the first\nTerminal Tab\n";\nkubectl proxy ↵
```

```
$ kubectl versionecho -e "\n\n\n\033[92mStarting Proxy. After starting it will not\nPlease click the first Terminal Tab\n";\n\nCommand 'kubectl' not found, did you mean:\n\n  command 'kubectl' from snap kubectl (1.20.4)\n\nSee 'snap info <snapname>' for additional versions.\n\n$ kubectl proxy\nStarting to serve on 127.0.0.1:8001\n\n
```

We now have a connection between our host (the online terminal) and the Kubernetes cluster. The proxy enables direct access to the API from these terminals.

You can see all those APIs hosted through the proxy endpoint. For example, we can query the version directly through the API using the `curl` command:

```
curl http://localhost:8001/version ↵
```

```
$ curl http://localhost:8001/version\n{\n    \"major\": \"1\", \n    \"minor\": \"20\", \n    \"gitVersion\": \"v1.20.2\", \n    \"gitCommit\": \"faecb196815e248d3ecfb03c680a4507229c2a56\", \n    \"gitTreeState\": \"clean\", \n    \"buildDate\": \"2021-01-13T13:20:00Z\", \n    \"goVersion\": \"go1.15.5\", \n    \"compiler\": \"gc\", \n    \"platform\": \"linux/amd64\"\n}\n\n
```

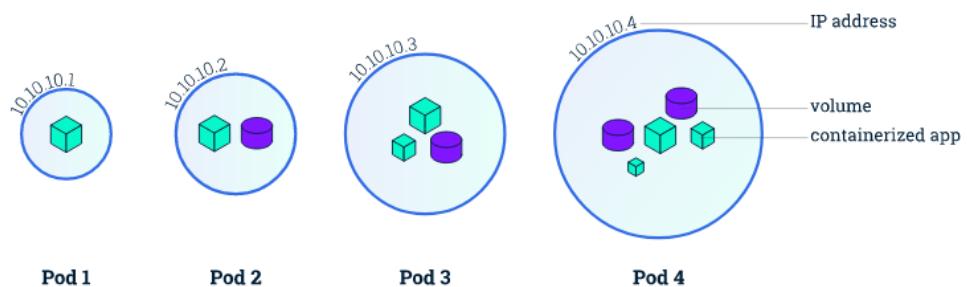
Viewing Pods and Nodes

Objectives

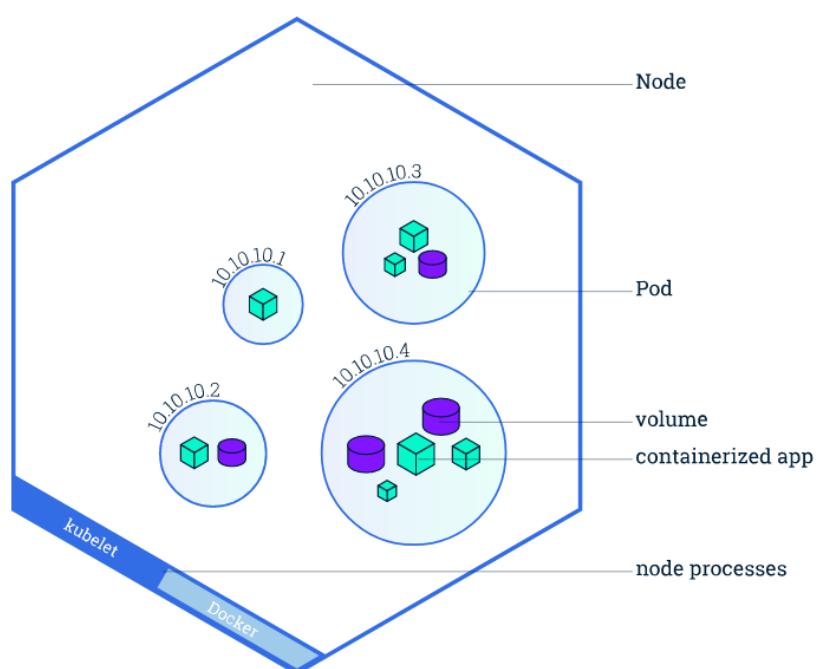
- Learn about Kubernetes Pods.
- Learn about Kubernetes Nodes

- Learn about Kubernetes nodes.
- Troubleshoot deployed applications.

Pods overview



Node overview



Troubleshooting with kubectl

In Module 2, you used Kubectl command-line interface. You'll continue to use it in Module 3 to get information about deployed applications and their environments. The most common operations can be done with the following kubectl commands:

- **kubectl get** - list resources
- **kubectl describe** - show detailed information about a resource
- **kubectl logs** - print the logs from a container in a pod
- **kubectl exec** - execute a command on a container in a pod

You can use these commands to see when applications were deployed, what their current statuses are, where they are running and what their configurations are.

Now that we know more about our cluster components and the command line, let's explore our application.

Step 1 Check application

configuration

Let's verify that the application we deployed in the previous scenario is running. We'll use the `kubectl get` command and look for existing Pods:

```
kubectl get pods ↵
```

If no pods are running then it means the interactive environment is still reloading it's previous state. Please wait a couple of seconds and list the Pods again. You can continue once you see the one Pod running.

```
NAME                      READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-fb5c67579-tr6rc   1/1     Running   0          56s
$ ↵
```

Next, to view what containers are inside that Pod and what images are used to build those containers we run the `describe pods` command:

```
kubectl describe pods ↵
```

We see here details about the Pod's container: IP address, the ports used and a list of events related to the lifecycle of the Pod.

The output of the `describe` command is extensive and covers some concepts that we didn't explain yet, but don't worry, they will become familiar by the end of this bootcamp.

```
$ ↵
$ kubectl get pods
NAME                      READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-fb5c67579-tr6rc   1/1     Running   0          56s
$ kubectl describe pods
Name:            kubernetes-bootcamp-fb5c67579-tr6rc
Namespace:       default
Priority:        0
Node:            minikube/10.0.0.16
Start Time:      Tue, 07 Feb 2023 09:27:15 +0000
Labels:          app=kubernetes-bootcamp
                 pod-template-hash=fb5c67579
Annotations:     <none>
Status:          Running
IP:              172.18.0.5
IPs:
  IP:           172.18.0.5
Controlled By:  ReplicaSet/kubernetes-bootcamp-fb5c67579
$ ↵
```

Using a Service to Expose Your App

Objectives

- Learn about a Service in Kubernetes
- Understand how labels and LabelSelector objects relate to a Service
- Expose an application outside a Kubernetes cluster using a Service

Overview of Kubernetes Services

Kubernetes [Pods](#) are mortal. Pods have a [lifecycle](#). When a worker node dies, the Pods running on the Node are also lost. A [ReplicaSet](#) might then dynamically drive the cluster back to the desired state via the creation of new Pods to keep your application running. As another example, consider an image-processing backend with 3 replicas. Those replicas are exchangeable; the front-end system should not care about backend replicas or even if a Pod is lost and recreated. That said, each Pod in a Kubernetes cluster has a unique IP address, even Pods on the same Node, so there needs to be a way of automatically reconciling changes among Pods so that your applications continue to function.

A Service in Kubernetes is an abstraction which defines a logical set of Pods and a policy by which to access them. Services enable a loose coupling between dependent Pods. A Service is defined using YAML ([preferred](#)) or JSON, like all Kubernetes objects. The set of Pods targeted by a Service is usually determined by a [LabelSelector](#) (see below for why you might want a Service without including a `selector` in the spec).

Although each Pod has a unique IP address, those IPs are not exposed outside the cluster without a Service. Services allow your applications to receive traffic. Services can be exposed in different ways by specifying a `type` in the `ServiceSpec`:

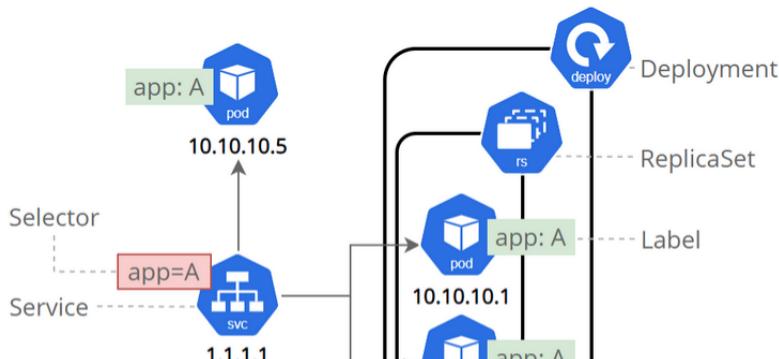
- *ClusterIP* (default) - Exposes the Service on an internal IP in the cluster. This type makes the Service only reachable from within the cluster.
- *NodePort* - Exposes the Service on the same port of each selected Node in the cluster using NAT. Makes a Service accessible from outside the cluster using `<NodeIP>:<NodePort>`. Superset of ClusterIP.
- *LoadBalancer* - Creates an external load balancer in the current cloud (if supported) and assigns a fixed, external IP to the Service. Superset of NodePort.
- *ExternalName* - Maps the Service to the contents of the `externalName` field (e.g. `foo.bar.example.com`), by returning a `CNAME` record with its value. No proxying of any kind is set up. This type requires v1.7 or higher of `kube-dns`, or CoreDNS version 0.0.8 or higher.

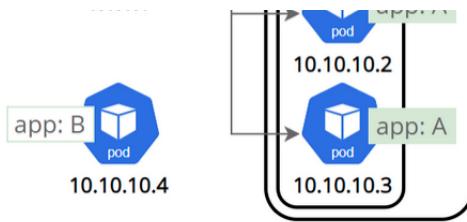
Services and Labels

A Service routes traffic across a set of Pods. Services are the abstraction that allows pods to die and replicate in Kubernetes without impacting your application. Discovery and routing among dependent Pods (such as the frontend and backend components in an application) are handled by Kubernetes Services.

Services match a set of Pods using [labels and selectors](#), a grouping primitive that allows logical operation on objects in Kubernetes. Labels are key/value pairs attached to objects and can be used in any number of ways:

- Designate objects for development, test, and production
- Embed version tags
- Classify an object using tags





Step 1 Create a new service

Let's verify that our application is running. We'll use the `kubectl get` command and look for existing Pods:

```
kubectl get pods ↵
```

If no pods are running then it means the interactive environment is still reloading its previous state. Please wait a couple of seconds and list the Pods again. You can continue once you see the one Pod running.

```
$ kubectl get pods
NAME                               READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-fb5c67579-ngrnh   1/1     Running   0          2m10s
$ kubectl get services
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes   ClusterIP   10.96.0.1      <none>        443/TCP      2m44s
```

```
kubectl get services ↵
```

We have a Service called `kubernetes` that is created by default when minikube starts the cluster. To create a new service and expose it to external traffic we'll use the `expose` command with `NodePort` as parameter (minikube does not support the `LoadBalancer` option yet).

```
$ kubectl expose deployment/kubernetes-bootcamp --type="NodePort" --port 8080
service/kubernetes-bootcamp exposed
$ kubectl get services
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes   ClusterIP   10.96.0.1      <none>        443/TCP      6m37s
```

```
kubernetes-bootcamp   NodePort    10.99.170.251   <none>        8080:31789/TCP   4
$
```

To find out what port was opened externally (by the NodePort option) we'll run the `describe service` command:

```
kubectl describe services/kubernetes-bootcamp ↵
```

Create an environment variable called `NODE_PORT` that has the value of the Node port assigned:

```
export NODE_PORT=$(kubectl get
services/kubernetes-bootcamp -o go-
template='{{(index .spec.ports
0).nodePort}}')
echo NODE_PORT=$NODE_PORT ↵
```

```
$ kubectl describe services/kubernetes-bootcamp
Name:                   kubernetes-bootcamp
Namespace:              default
Labels:                 app=kubernetes-bootcamp
Annotations:            <none>
Selector:               app=kubernetes-bootcamp
Type:                  NodePort
IP Families:           <none>
IP:                    10.104.173.224
IPs:                  10.104.173.224
Port:                  <unset>  8080/TCP
TargetPort:             8080/TCP
NodePort:               <unset>  31257/TCP
Endpoints:              172.18.0.5:8080
Session Affinity:       None
External Traffic Policy: Cluster
Events:                <none>
$ ↵
```

```
$ export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
$ echo NODE_PORT=$NODE_PORT
NODE_PORT=31257
$ ↵
```

```
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-fb5c67579-4dbrt | v=1
$ ↵
```

Now we can test that the app is exposed outside of the cluster using `curl`, the IP of the Node and the externally exposed port:

```
curl $(minikube ip):$NODE_PORT ↵
```

And we get a response from the server. The Service is exposed.

Step 2: Using labels

The Deployment created automatically a label for our Pod. With `describe deployment` command you can see the name of the label:

```
kubectl describe deployment ✓
```

Let's use this label to query our list of Pods. We'll use the `kubectl get pods` command with `-l` as a parameter, followed by the label values:

```
kubectl get pods -l app=kubernetes-bootcamp ✓
```

You can do the same to list the existing services:

```
kubectl get services -l app=kubernetes-bootcamp ✓
```

```
$ kubectl get pods -l app=kubernetes-bootcamp
NAME                               READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-fb5c67579-4dbrt   1/1     Running   0          4m51s
$ kubectl get services -l app=kubernetes-bootcamp
NAME           TYPE      CLUSTER-IP        EXTERNAL-IP      PORT(S)         AGE
kubernetes-bootcamp   NodePort   10.104.173.224   <none>        8080:31257/TCP   5m7s
$ [ ]
```

Get the name of the Pod and store it in the `POD_NAME` environment variable:

```
export POD_NAME=$(kubectl get pods -o
go-template --template '{{range
.items}}{{.metadata.name}}\n{{end}}')
echo Name of the Pod: $POD_NAME ✓
```

To apply a new label we use the `label` command followed by the object type, object name and the new label:

```
kubectl label pods $POD_NAME  
version=v1 ↵
```

```
$ export POD_NAME=$(kubectl get pods -o go-template --template '{{range .items}}{{.metadata.name}}\n{{end}}')  
$ echo Name of the Pod: $POD_NAME  
Name of the Pod: kubernetes-bootcamp-fb5c67579-4dbrt  
$ kubectl label pods $POD_NAME version=v1  
pod/kubernetes-bootcamp-fb5c67579-4dbrt labeled  
$ ↵
```

This will apply a new label to our Pod (we pinned the application version to the Pod), and we can check it with the describe pod command:

```
kubectl describe pods $POD_NAME ↵
```

We see here that the label is attached now to our Pod. And we can query now the list of pods using the new label:

```
kubectl get pods -l version=v1 ↵
```

And we see the Pod.

```
Port:          8080/TCP  
$ kubectl describe pods $POD_NAME  
Name:          kubernetes-bootcamp-fb5c67579-4dbrt  
Namespace:     default  
Priority:      0  
Node:          minikube/10.0.0.6  
Start Time:    Tue, 07 Feb 2023 10:24:53 +0000  
Labels:        app=kubernetes-bootcamp  
               pod-template-hash=fb5c67579  
               version=v1  
Annotations:   <none>  
Status:        Running  
IP:            172.18.0.5  
IPs:  
  IP:          172.18.0.5  
Controlled By: ReplicaSet/kubernetes-bootcamp-fb5c67579  
Containers:    ↵
```

```
$ kubectl get pods -l version=v1  
NAME                  READY   STATUS    RESTARTS   AGE  
kubernetes-bootcamp-fb5c67579-4dbrt   1/1     Running   0          9m15s  
$ ↵
```

Step 3 Deleting a

Step 5: Deleting a service

To delete Services you can use the `delete service` command. Labels can be used also here:

```
kubectl delete service -l  
app=kubernetes-bootcamp ↵
```

Confirm that the service is gone:

```
kubectl get services ↵
```

```
$ kubectl delete service -l app=kubernetes-bootcamp  
service "kubernetes-bootcamp" deleted  
$ kubectl get services  
NAME          TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)    AGE  
kubernetes    ClusterIP  10.96.0.1    <none>        443/TCP   11m  
$ ↵
```

This confirms that our Service was removed. To confirm that route is not exposed anymore you can `curl` the previously exposed IP and port:

```
curl $(minikube ip):$NODE_PORT ↵
```

This proves that the app is not reachable anymore from outside of the cluster. You can confirm that the app is still running with a `curl` inside the pod:

```
kubectl exec -ti $POD_NAME -- curl  
localhost:8080 ↵
```

We see here that the application is up. This is because the Deployment is managing the application. To shut down the application, you would need to delete the Deployment as well.

```
$ curl $(minikube ip):$NODE_PORT
curl: (7) Failed to connect to 10.0.0.6 port 31257: Connection refused
$ kubectl exec -ti $POD_NAME -- curl localhost:8080
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-fb5c67579-4dbrt | v=1
$ 
```

Installation :

<https://kubernetes.io/fr/docs/setup/>

<https://minikube.sigs.k8s.io/docs/start/>

1 Installation

Click on the buttons that describe your target platform. For other architectures, see the [release page](#) for a complete list of minikube binaries.

Operating system	<input checked="" type="button"/> Linux	<input type="button"/> macOS	<input type="button"/> Windows		
Architecture	<input checked="" type="button"/> x86-64	<input type="button"/> ARM64	<input type="button"/> ARMv7	<input type="button"/> ppc64	<input type="button"/> S390x
Release type	<input checked="" type="button"/> Stable	<input type="button"/> Beta			
Installer type	<input type="button"/> Binary download	<input type="button"/> Debian package	<input type="button"/> RPM package		

To install the latest minikube **stable** release on **x86-64 Linux** using **binary download**:

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

```
vel@vel-Precision-3650-Tower:~$ curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
% Total    % Received % Xferd  Average Speed   Time      Time     Current
          Dload  Upload Total   Spent    Left Speed
100  77.3M  100  77.3M    0     0  64.5M      0  0:00:01  0:00:01  --:--:-- 64.6M
vel@vel-Precision-3650-Tower:~$ sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

2 Start your cluster

From a terminal with administrator access (but not logged in as root), run:

```
minikube start
```

If minikube fails to start, see the [drivers page](#) for help setting up a compatible container or virtual-machine manager.

```
vel@vel-Precision-3650-Tower:~$ minikube start
minikube v1.29.0 sur Ubuntu 22.04
Choix automatique du pilote docker. Autres choix: qemu2, virtualbox, ssh, none
Utilisation du pilote Docker avec le privilège root
Démarrage du noeud de plan de contrôle minikube dans le cluster minikube
Extraction de l'image de base...
Téléchargement du préchargement de Kubernetes v1.26.1...
> preloaded-images-k8s-v18-v1...: 397.05 MiB / 397.05 MiB 100.00% 60.00 M
> gcr.io/k8s-minikube/kicbase...: 407.19 MiB / 407.19 MiB 100.00% 32.34 M
Création de docker container (CPUs=2, Memory=7900Mo) ...
Préparation de Kubernetes v1.26.1 sur Docker 20.10.23...
■ Génération des certificats et des clés
■ Démarrage du plan de contrôle ...
■ Configuration des règles RBAC ...
Configuration de bridge CNI (Container Networking Interface)...
■ Utilisation de l'image gcr.io/k8s-minikube/storage-provisioner:v5
Modules activés: storage-provisioner, default-storageclass
Vérification des composants Kubernetes...
kubectl introuvable. Si vous en avez besoin, essayez : 'minikube kubectl -- get pods -A'
Terminé ! kubectl est maintenant configuré pour utiliser "minikube" cluster et espace de noms "default" par défaut.
```

3 Interact with your cluster

If you already have kubectl installed, you can now use it to access your shiny new cluster:

```
kubectl get po -A
```

Alternatively, minikube can download the appropriate version of kubectl and you should be able to use it like this:

```
minikube kubectl -- get po -A
```

You can also make your life easier by adding the following to your shell config:

```
alias kubectl="minikube kubectl --"
```

Initially, some services such as the storage-provisioner, may not yet be in a Running state. This is a normal condition during cluster bring-up, and will resolve itself momentarily. For additional insight into your cluster state, minikube bundles the Kubernetes Dashboard, allowing you to get easily acclimated to your new environment:

```
minikube dashboard
```

Pour allez plus loin dans le tutoriel, il faut installer kubectl

Installer Kubectl (deux façon de faire : curl et snap) :

<https://kubernetes.io/docs/tasks/tools/install-kubectl-linux/>

Install kubectl binary with curl on Linux

1. Download the latest release with the command:

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
```

Note:

To download a specific version, replace the `$(curl -L -s https://dl.k8s.io/release/stable.txt)` portion of the command with the specific version.

For example, to download version v1.26.0 on Linux, type:

```
curl -LO https://dl.k8s.io/release/v1.26.0/bin/linux/amd64/kubectl
```

2. Validate the binary (optional)

Download the kubectl checksum file:

```
curl -LO "https://dl.k8s.io/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl.sha256"
```

Validate the kubectl binary against the checksum file:

```
echo "$(cat kubectl.sha256) kubectl" | sha256sum --check
```

If valid, the output is:

```
kubectl: OK
```

If the check fails, `sha256` exits with nonzero status and prints output similar to:

```
kubectl: FAILED
sha256sum: WARNING: 1 computed checksum did NOT match
```

Note: Download the same version of the binary and checksum.

3. Install kubectl

```
sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
```

4. Test to ensure the version you installed is up-to-date:

```
kubectl version --client
```

```
vel@vel-Precision-3650-Tower:~$ curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
% Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload   Total Spent  Left  Speed
100 138 100 138   0     0  726      0 --::-- --::-- --::--  730
100 45.7M 100 45.7M  0     0  46.3M      0 --::-- --::-- --::-- 100M
vel@vel-Precision-3650-Tower:~$ curl -LO "https://dl.k8s.io/S(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl.sha256"
% Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload   Total Spent  Left  Speed
100 138 100 138   0     0  756      0 --::-- --::-- --::--  754
100  64 100  64   0     0  181      0 --::-- --::-- --::-- 181
vel@vel-Precision-3650-Tower:~$ echo "$(
```

Autre façon d'installer kubectl : snap

```
snap install kubectl --classic
kubectl version --client
```

```
vel@vel-Precision-3650-Tower:~$ snap install kubectl --classic
kubectl 1.26.1 par Canonical✓ installé
vel@vel-Precision-3650-Tower:~$ kubectl version --client
WARNING: This version information is deprecated and will be replaced with the output from kubectl version --short. Use --output=yaml|json to get the full version.
Client Version: version.Info{Major:"1", Minor:"26", GitVersion:"v1.26.1", GitCommit:"8f94681cd294aa8cf3407b8191f6c70214973a4", GitTreeState:"clean", BuildDate:"2023-01-18T15:58:16Z", GoVersion:"go1.19.5", Compiler:"gc", Platform:"linux/amd64"}
Kustomize Version: v4.5.7
```

Une fois la commande kubectl installé : on peut reprendre l'initiation à Minikube :

Alternatively, minikube can download the appro

```
minikube kubectl -- get po -A
```

You can also make your life easier by adding the

```
alias kubectl="minikube kubectl --"
```

Initially, some services such as the storage-provider will bring-up, and will resolve itself moment Kubernetes Dashboard, allowing you to get easily

```
minikube dashboard
```

```
vel@vel-Precision-3650-Tower:~$ kubectl get po -A
NAMESPACE     NAME           READY   STATUS    RESTARTS   AGE
kube-system   coredns-787d4945fb-nzfh9   1/1     Running   0          12m
kube-system   etcd-minikube        1/1     Running   0          12m
kube-system   kube-apiserver-minikube  1/1     Running   0          12m
```

kube-system	kube-controller-manager-minikube	1/1	Running	0	12m
kube-system	kube-proxy-trfzn	1/1	Running	0	12m
kube-system	kube-scheduler-minikube	1/1	Running	0	12m
kube-system	storage-provisioner	1/1	Running	1 (12m ago)	12m

Pour ce tp, je ne ferais pas l'alias

Déploiement du dashboard :

```
vel@vel-Precision-3650-Tower:~$ minikube dashboard
👉 Activation du tableau de bord...
  └─ Utilisation de l'image docker.io/kubernetesui/dashboard:v2.7.0
  └─ Utilisation de l'image docker.io/kubernetesui/metrics-scraper:v1.0.8
💡 Certaines fonctionnalités du tableau de bord nécessitent le module metrics-server. Pour activer toutes les fonctionnalités, veuillez exécuter :
  minikube addons enable metrics-server

⚠️ Vérification de l'état du tableau de bord...
  └─ Lancement du proxy...
  └─ Vérification de l'état du proxy...
  └─ Ouverture de http://127.0.0.1:33073/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:/proxy/ dans votre navigateur par défaut..
gtk-Message: 12:03:45.257: Not loading module "atk-bridge": The functionality is provided by GTK natively. Please try to not load it.
└─
```

Vérification des containers : on s'aperçoit que notre dashboard s'exécute via dockers

```
vel@vel-Precision-3650-Tower:~$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
MES
4d089f0201a2      gcr.io/k8s-minikube/kicbase:v0.0.37   "/usr/local/bin/entr..."   21 minutes ago    Up 21 minutes   127.0.0.1:49151->22/tcp, 127.0.0.1:49156->2376/tcp, 127.0.0.1:49155->5000/tcp, 127.0.0.1:49154->8443/tcp, 127.0.0.1:49153->32443/tcp
57->22/tcp, 127.0.0.1:49156->2376/tcp, 127.0.0.1:49155->5000/tcp, 127.0.0.1:49154->8443/tcp, 127.0.0.1:49153->32443/tcp
minikube
36011283a72a      portainer/portainer-ce:2.9.3       "/portainer"          5 days ago       Up 3 hours      0.0.0.0:8000->8000/tcp, :::8000->8000/tcp, 0.0.0.0:9443->9443/tcp, :::9443->9443/tcp, 9000/tcp
rtainer
```

Déploiement application :

4 Deploy applications

Service LoadBalancer Ingress

Create a sample deployment and expose it on port 8080:

```
kubectl create deployment hello-minikube --image=kicbase/echo-server:1.0
kubectl expose deployment hello-minikube --type=NodePort --port=8080
```

It may take a moment, but your deployment will soon show up when you run:

```
kubectl get services hello-minikube
```

The easiest way to access this service is to let minikube launch a web browser for you:

```
minikube service hello-minikube
Service
```

Alternatively, use kubectl to forward the port:

```
kubectl port-forward service/hello-minikube 7080:8080
```

Tada! Your application is now available at <http://localhost:7080/>.

You should be able to see the request metadata in the application output. Try changing the path of the request and observe the changes. Similarly, you can do a POST request and observe the body show up in the output.

On déploie une application qu'on expose sur le port 8080

```
vel@vel-Precision-3650-Tower:~$ kubectl create deployment hello-minikube --image=kicbase/echo-server:1.0
deployment.apps/hello-minikube created
vel@vel-Precision-3650-Tower:~$ kubectl expose deployment hello-minikube --type=NodePort --port=8080
service/hello-minikube exposed
vel@vel-Precision-3650-Tower:~$
```

Vérification des services lancés

```
vel@vel-Precision-3650-Tower:~$ kubectl get services hello-minikube
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
hello-minikube  NodePort  10.100.22.94 <none>       8080:31193/TCP 64s
```

```
vel@vel-Precision-3650-Tower:~$ minikube service hello-minikube
|-----|-----|-----|-----|
| NAMESPACE | NAME | TARGET PORT | URL |
|-----|-----|-----|-----|
| default | hello-minikube | 8080 | http://192.168.49.2:31193 |
|-----|-----|-----|-----|
💡 Ouverture du service default/hello-minikube dans le navigateur par défaut...
vel@vel-Precision-3650-Tower:~$ Gtk-Message: 12:23:42.857: Not loading module "atk-bridge": The functionality is provided by GTK natively. Please try to not load it.
```

```
Request served by hello-minikube-77b6f68484-68w67
HTTP/1.1 GET /
Host: 192.168.49.2:31193
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*/*;q=0.8
Accept-Encoding: gzip, deflate
Accept-Language: fr,fr-FR;q=0.8,en-US;q=0.5,en;q=0.3
Connection: keep-alive
Upgrade-Insecure-Requests: 1
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:109.0) Gecko/20100101 Firefox/109.0
```

On constate également le rajout du pods sur la page de management :

Workloads	Nombre	Etat
Cron Jobs	0	
Daemon Sets	0	
Deployments	1	Running 1
Jobs	0	
Pods	1	Running 1
Replica Sets	0	
Replication Controllers	0	
Stateful Sets	0	

Service	Nombre	Etat
Ingresses	0	
Ingress Classes	0	
Services	1	Running 1

Config and Storage	Nombre	Etat
Config Maps	0	
Persistent Volume Claims	0	

On utilise maintenant kubectl pour ouvrir les ports

```
vel@vel-Precision-3650-Tower:~$ kubectl port-forward service/hello-minikube 7080:8080
Forwarding from 127.0.0.1:7080 -> 8080
Forwarding from [::1]:7080 -> 8080
```

```
Request served by hello-minikube-77b6f68484-68w67
HTTP/1.1 GET /
Host: localhost:7080
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*/*;q=0.8
Accept-Encoding: gzip, deflate, br
Accept-Language: fr,fr-FR;q=0.8,en-US;q=0.5,en;q=0.3
Connection: keep-alive
Cookie: jenkins-timestamper-offset=-3600000
Sec-Fetch-Dest: document
Sec-Fetch-Mode: navigate
```

```
Sec-Fetch-Site: cross-site
Sec-Fetch-User: ?1
Upgrade-Insecure-Requests: 1
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:109.0) Gecko/20100101 Firefox/109.0
```

Service LoadBalancer Ingress

To access a LoadBalancer deployment, use the “minikube tunnel” command. Here is an example deployment:

```
kubectl create deployment balanced --image=kicbase/echo-server:1.0
kubectl expose deployment balanced --type=LoadBalancer --port=8080
```

In another window, start the tunnel to create a routable IP for the ‘balanced’ deployment:

```
minikube tunnel
```

To find the routable IP, run this command and examine the `EXTERNAL-IP` column:

```
kubectl get services balanced
```

Your deployment is now available at <EXTERNAL-IP>:8080

```
vel@vel-Precision-3650-Tower:~$ kubectl create deployment balanced --image=kicbase/echo-server:1.0
deployment.apps/balanced created
vel@vel-Precision-3650-Tower:~$ kubectl expose deployment balanced --type=LoadBalancer --port=8080
service/balanced exposed
```

```
vel@vel-Precision-3650-Tower:~$ minikube tunnel
[sudo] Mot de passe de vel :
Status:
  machine: minikube
  pid: 67575
  route: 10.96.0.0/12 -> 192.168.49.2
  minikube: Running
  services: [balanced]
errors:
  minikube: no errors
  router: no errors
  loadbalancer emulator: no errors
```

Déploiements				
Nom	Images	Étiquettes		
balanced	kicbase/echo-server:1.0	app: balanced		
hello-minikube	kicbase/echo-server:1.0	app: hello-minikube		
Nom	Images	Étiquettes	Noeud	Statut
balanced-56c586fb9-d7bt4	kicbase/echo-server:1.0	app: balanced pod-template-hash: 56c586fb9	minikube	Running
hello-minikube-77b6f6848-68w67	kicbase/echo-server:1.0	app: hello-minikube pod-template-hash: 77b6f6848	minikube	Running

```
vel@vel-Precision-3650-Tower:~$ kubectl get services balanced
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
balanced   LoadBalancer   10.101.53.141   10.101.53.141   8080/TCP   117s
```

```
balanced LoadBalancer 10.100.53.141 10.100.53.141 8080:30182/TCP 117s
vel@vel-Precision-3650-Tower:~$
```

Addon ingress :

Service LoadBalancer **Ingress**

Enable ingress addon:

```
minikube addons enable ingress
```

The following example creates simple echo-server services and an Ingress object to route to these services.

```
vel@vel-Precision-3650-Tower:~$ minikube addons enable ingress
💡 ingress est un addon maintenu par Kubernetes. Pour toute question, contactez minikube sur GitHub.
Vous pouvez consulter la liste des mainteneurs de minikube sur : https://github.com/kubernetes/minikube/blob/master/OWNERS
  ■ Utilisation de l'image registry.k8s.io/ingress-nginx/kube-webhook-certgen:v20220916-gd32f8c343
  ■ Utilisation de l'image registry.k8s.io/ingress-nginx/kube-webhook-certgen:v20220916-gd32f8c343
  ■ Utilisation de l'image registry.k8s.io/ingress-nginx/controller:v1.5.1
  🔍 Vérification du module ingress...
🌟 Le module 'ingress' est activé
vel@vel-Precision-3650-Tower:~$
```

Apply the contents

```
kubectl apply -f https://storage.googleapis.com/minikube-site-examples/ingress-example.yaml
```

Wait for ingress address

```
kubectl get ingress
NAME      CLASS    HOSTS    ADDRESS        PORTS   AGE
example-ingress  nginx  *     <your_ip_here>  80      5m45s
```

Note for Docker Desktop Users:

To get ingress to work you'll need to open a new terminal window and run `minikube tunnel` and in the following step use `127.0.0.1` in place of `<ip_from_above>`.

Now verify that the ingress works

```
$ curl <ip_from_above>/foo
Request served by foo-app
...
$ curl <ip_from_above>/bar
Request served by bar-app
...
```

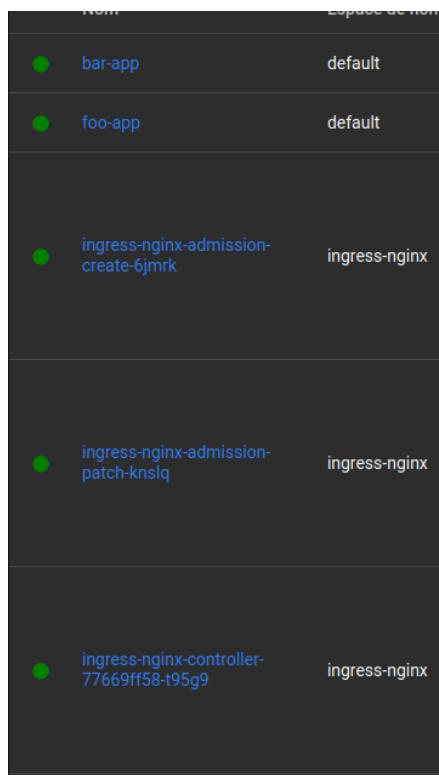
```
vel@vel-Precision-3650-Tower:~$ minikube addons enable ingress
💡 ingress est un addon maintenu par Kubernetes. Pour toute question, contactez minikube sur : https://github.com/kubernetes/minikube/blob/master/OWNERS
  ■ Utilisation de l'image registry.k8s.io/ingress-nginx/kube-webhook-certgen:v20220916-gd32f8c343
  ■ Utilisation de l'image registry.k8s.io/ingress-nginx/kube-webhook-certgen:v20220916-gd32f8c343
  ■ Utilisation de l'image registry.k8s.io/ingress-nginx/controller:v1.5.1
  🔍 Vérification du module ingress...
🌟 Le module 'ingress' est activé
vel@vel-Precision-3650-Tower:~$ kubectl apply -f https://storage.googleapis.com/minikube-site-examples/ingress-example.yaml
pod/foo-app created
service/foo-service created
pod/bar-app created
service/bar-service created
ingress.networking.k8s.io/example-ingress created
vel@vel-Precision-3650-Tower:~$ kubectl get ingress
NAME      CLASS    HOSTS    ADDRESS        PORTS   AGE
example-ingress  nginx  *     192.168.49.2  80      85s
vel@vel-Precision-3650-Tower:~$ curl 192.168.49.2/foo
Request served by foo-app
HTTP/1.1 GET /foo
Host: 192.168.49.2
```

```

HOST: 192.168.49.2
Accept: /*
User-Agent: curl/7.81.0
X-Forwarded-For: 192.168.49.1
X-Forwarded-Host: 192.168.49.2
X-Forwarded-Port: 80
X-Forwarded-Proto: http
X-Forwarded-Scheme: http
X-Real-Ip: 192.168.49.1
X-Request-Id: 170a97423fb3cc90e21018ae5422e247
X-Scheme: http
vel@vel-Precision-3650-Tower:~$ curl 192.168.49.2/bar

```

Suite à ces manipulations, on voit que nous avons 5 nouveaux pods :



Suite du TP : tenter de lancer une image nginx avec minikube :

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80

```

Minicube sert uniquement à créer des VM en local, il ne nous sera pas utilisé pour ce TP

Pour faire ce TP, il faudrait utiliser KubeADM :

<https://kubernetes.io/fr/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>

Installation KubeADM

Vérifiez les ports requis

noeuds maîtres (masters)

Protocole	Direction	Plage de Port	Utilisé pour	Utilisé par
-----------	-----------	---------------	--------------	-------------

TCP	Entrant	6443*	Kubernetes API server	Tous
TCP	Entrant	2379-2380	Etcd server client API	kube-apiserver, etcd
TCP	Entrant	10250	Kubelet API	Lui-même, Control plane
TCP	Entrant	10251	kube-scheduler	Lui-même
TCP	Entrant	10252	kube-controller-manager	Lui-même

nœuds workers

Protocole	Direction	Plage de Port	Utilisé pour	Utilisé par
TCP	Entrant	10250	Kubelet API	Lui-même, Control plane
TCP	Entrant	30000-32767	NodePort Services**	Eux-mêmes

** Plage de ports par défaut pour les [Services NodePort](#).

Tous les numéros de port marqués d'un * sont écrasables. Vous devrez donc vous assurer que les ports personnalisés que vous utilisez sont également ouverts.

Bien que les ports etcd soient inclus dans les nœuds masters, vous pouvez également héberger votre propre cluster etcd en externe ou sur des ports personnalisés.

Le plug-in de réseau de pod que vous utilisez (voir ci-dessous) peut également nécessiter certains ports à ouvrir. Étant donné que cela diffère d'un plugin à l'autre, veuillez vous reporter à la documentation des plugins sur le(s) port(s) requis(s).

Par défaut, Kubernetes utilise le [Container Runtime Interface \(CRI\)](#) pour s'interfacer avec votre environnement d'exécution de conteneur choisi.

Si vous ne spécifiez pas de runtime, kubeadm essaie automatiquement de détecter un Runtime de conteneur en parcourant une liste de sockets de domaine Unix bien connus. Le tableau suivant répertorie les environnements d'exécution des conteneurs et leurs chemins de socket associés:

Runtime Chemin vers le socket de domaine Unix

Docker	/var/run/docker.sock
containerd	/run/containerd/containerd.sock
CRI-O	/var/run/crio/crio.sock

Si Docker et containerd sont détectés, Docker est prioritaire. C'est nécessaire car Docker 18.09 est livré avec containerd et les deux sont détectables même si vous installez Docker. Si deux autres environnements d'exécution ou plus sont détectés, kubeadm se ferme avec une erreur.

Le kubelet s'intègre à Docker via l'implémentation CRI intégrée de `dockershim`.

Voir [runtimes de conteneur](#) pour plus d'informations.

Installation de kubeadm, des kubelets et de kubectl

Vous installerez ces paquets sur toutes vos machines:

- `kubeadm` : la commande pour initialiser le cluster.
- la `kubelet` : le composant qui s'exécute sur toutes les machines de votre cluster et fait des actions comme le démarrage des pods et des conteneurs.
- `kubectl` : la ligne de commande utilisée pour parler à votre cluster.

kubeadm **n'installera pas** ni ne gérera les `kubelet` ou `kubectl` pour vous. Vous devez vous assurer qu'ils correspondent à la version du control plane de Kubernetes que vous souhaitez que kubeadm installe pour vous. Si vous ne le faites pas, vous risquez qu'une erreur de version se produise, qui pourrait conduire à un comportement inattendu. Cependant, une version mineure entre les kubelets et le control plane est pris en charge, mais la version de la kubelet ne doit jamais dépasser la version de l'API server. Par exemple, les kubelets exécutant la version 1.7.0 devraient être entièrement compatibles avec un API server en 1.8.0, mais pas l'inverse.

For information about installing `kubectl`, see [Installation et configuration kubectl](#).

Attention: Ces instructions excluent tous les packages Kubernetes de toutes les mises à niveau du système d'exploitation. C'est parce que kubeadm et Kubernetes ont besoin d'une [attention particulière lors de la mise à niveau](#).

Ubuntu, Debian or HypriotOS

CentOS, RHEL or Fedora

Fedora CoreOS ou Flatcar Container Linux

```
sudo apt-get update && sudo apt-get install -y apt-transport-https curl
curl -s https://nackanes.cloud.google.com/ant/doc/ant-key.gpg | sudo ant-key add -
```

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add
cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF
sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo apt-mark hold kubelet kubeadm kubectl
```

Kubelet redémarre maintenant toutes les quelques secondes, car il attend les instructions de kubeadm dans une boucle de crash.

```
● vel@vel-Precision-3650-Tower:~$ curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).
OK
● vel@vel-Precision-3650-Tower:~$ cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
> deb https://apt.kubernetes.io/ kubernetes-xenial main
> EOF
deb https://apt.kubernetes.io/ kubernetes-xenial main
○ vel@vel-Precision-3650-Tower:~$ sudo apt-get update
Atteint :2 https://ppa.launchpadcontent.net/ansible/ubuntu jammy InRelease
Atteint :3 https://ppa.launchpadcontent.net/openjdk-r/ppa/ubuntu jammy InRelease
Réception de :1 https://packages.cloud.google.com/apt kubernetes-xenial InRelease [8 993 B]
Réception de :4 https://ngrok-agent.s3.amazonaws.com buster InRelease [20,3 kB]
Réception de :5 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 Packages [63,2 kB]
```

Installation KubeADM sur Arch :

<https://wiki.archlinux.org/title/Kubernetes>

1 Installation

When manually creating a Kubernetes cluster **install etcd^{AUR}** and the **package group kubernetes-control-plane** (for a control-plane node) and **kubernetes-node** (for a worker node).

When creating a Kubernetes cluster with the help of **kubeadm**, **install kubeadm** and **kubelet** on each node.

Both control-plane and regular worker nodes require a container runtime for their **kubelet** instances which is used for hosting containers. **Install either containerd or cri-o** to meet this dependency.

To control a kubernetes cluster, **install kubectl** on the control-plane hosts and any external host that is supposed to be able to interact with the cluster.

helm is a tool for managing pre-configured Kubernetes resources which may be helpful for getting started.

Objectifs

- Installer un cluster Kubernetes à master unique ou un [cluster à haute disponibilité](#)
- Installez un réseau de pods sur le cluster afin que vos pods puissent se parler

Initialiser votre master

Le master est la machine sur laquelle s'exécutent les composants du control plane, y compris etcd (la base de données du cluster) et l'API serveur (avec lequel la CLI kubectl communique).

1. Choisissez un add-on réseau pour les pods et vérifiez s'il nécessite des arguments à passer à l'initialisation de kubeadm. Selon le fournisseur tiers que vous choisissez, vous devrez peut-être définir le `--pod-network-cidr` sur une valeur spécifique au fournisseur. Voir [Installation d'un add-on réseau de pod](#).
2. (Facultatif) Sauf indication contraire, kubeadm utilise l'interface réseau associée avec la passerelle par défaut pour annoncer l'IP du master. Pour utiliser une autre interface réseau, spécifiez l'option `--apiserver-advertise-address=<ip-address>` à `kubeadm init`. Pour déployer un cluster Kubernetes en utilisant l'adressage IPv6, vous devez spécifier une adresse IPv6, par exemple `--apiserver-advertise-address=fd00::101`
3. (Optional) Lancez `kubeadm config images pull` avant de faire `kubeadm init` pour vérifier la connectivité aux registres gcr.io.

